




REINFORCEMENT LEARNING

JALEY DHOLAKIYA

RECAP : DEEP LEARNING




$$y \approx f_w(x)$$

PROGRAMMING VS DEEP LEARNING




PROGRAMMING



DEEP LEARNING

MNIST DIGIT CLASSIFICATION




$$y \approx f_w(x)$$

CHALLENGES IN DEEP LEARNING

- ITS DATA HUNGRY
- NO ACTIVE PARTICIPATION

REINFORCEMENT LEARNING



RL VS DL

- ACTIVE ENGAGEMENT
- GET RID OF MANUAL DATA COLLECTION

Components of RL



AGENT



POLICY = ACTION | STATE



STATE



OBSERVATION



ACTION



REWARD

ENVIRONMENT


- EPISODIC VS NON-EPISODIC
- FIXED VS DYNAMIC
- SINGLE VS MULTIPLAYER

VIDEO 1 : COMPONENTS OF RL

GYM EXAMPLE

GOAL OF RL

STATE	ACTION
A	GO STRAIGHT
B	TAKE LEFT
C	TAKE RIGHT OR LEFT




WHICH ACTION IS BEST?

MOST REWARDING ACTION

IMMEDIATE REWARD?



IMMEDIATE REWARD?



DISCOUNTED REWARD



TYPES OF RL ALGORITHMS

(BASED ON MODEL)

- MODEL DEPENDENT
(TOO MUCH MATH/ LESS USEFUL)
- MODEL FREE
(LESS MATH + MORE FUN)

TYPES OF RL ALGORITHMS

(BASED ON POLICY)

- POLICY EVALUATION


ESTIMATE THE STATE/ACTION REWARDS
FOR A GIVEN POLICY

- POLICY CONTROL


TASK IS TO IMPROVE A POLICY

POLICY CONTROL

EPSILON GREEDY POLICY




EXPLORE EXPLOIT PROBLEM



EXPLORE EXPLOIT PROBLEM

STATE VISITATION FREQUENCY



MONTÉ CARLO EVALUATION

Farm → Sea → Island → Forest → Dead

Sea → Island → Forest → Mountain → Forest → Dead

Forest → Sea → Island → Sea → Island → Forest → Dead



Monte Carlo Evaluation

Algorithm 1: Monte Carlo Evaluation

```
1  $N(s) \leftarrow 0, \forall s \in S$ 
2  $S(s) \leftarrow 0, \forall s \in S$ 
3 for each episode do
4   for each state( $s$ )  $\in$  episode do
5      $G_t \leftarrow 0$ 
6      $i = 0$ 
7     for  $t \leftarrow [t_{curr}, T]$  do
8        $G_t \leftarrow G_t + \gamma^i R_t$ 
9        $i \leftarrow i + 1$ 
10     $N(s) \leftarrow N(s) + 1$ 
11     $S(s) \leftarrow S(s) + G_t$ 
12     $V(s) \leftarrow S(s)/N(s)$ 
```

Algorithm 2: Monte Carlo Evaluation with incremental updates

```
1  $N(s) \leftarrow 0, \forall s \in S$ 
2  $S(s) \leftarrow 0, \forall s \in S$ 
3 for each episode do
4   for each state( $s$ )  $\in$  episode do
5      $G_t \leftarrow 0$ 
6      $i = 0$ 
7     for  $t \leftarrow [t_{curr}, T]$  do
8        $G_t \leftarrow G_t + \gamma^i R_t$ 
9        $i \leftarrow i + 1$ 
10     $V(s) \leftarrow V(s) + \alpha(G_t - V(s))$ 
11    or
12     $N(s) \leftarrow N(s) + 1$ 
13     $V(s) \leftarrow V(s) + \frac{1}{N(s)}(G_t - V(s))$ 
```

TEMPORAL DIFFERENCE

$$G_t(s_t) = R_{t+1} + \gamma R_{t+2} + \gamma^{T-2} \dots R_{T-1}$$

We can write reward for next state as

$$G_{t+1}(s_{t+1}) = R_{t+2} + \gamma R_{t+3} + \gamma^{T-3} \dots R_{T-1}$$

Therefore we can also write the first equation as

$$G_t(s_t) = R_{t+1} + \gamma \overbrace{\{R_{t+2} + \gamma^{T-3} \dots R_{T-1}\}}^{G_{t+1}(s_{t+1})}$$


and we all know we want to update value function, closer to discounted reward ($G \leftarrow V$), therefore our new update for TD Learning becomes

$$V(s_t) = V(s) + \alpha(\hat{G}_t - V(s))$$

$$V(s_t) = V(s_t) + \alpha(R_{t+1} + \gamma V(s_{t+1}) - V(s_t))$$


TEMPORAL DIFFERENCE

UPDATE $V(s_2)$ BASED ON R_3 and $V(s_3)$



REMINDER : *GOAL* OF RL


STATE	ACTION
A	GO STRAIGHT
B	TAKE LEFT
C	TAKE RIGHT OR LEFT



MC VS TD

MONTE CARLO


$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$



$$q(a_t|s_t) = q(a_t|s_t) + \alpha (G_t - q(a_t|s_t))$$

TEMPORAL DIFFERENCE

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



$$q(a_t|s_t) = q(a_t|s_t) + \alpha ((R_t + \gamma q(a_{t+1}|s_{t+1})) - q(a_t|s_t))$$

TD CONTROL + GREEDY = SARSA

(STATE ACTION REWARD STATE ACTION)

POLICY EVALUATION

$$q(a_t|s_t) = q(a_t|s_t) + \alpha((R_t + \gamma q(a_{t+1}|s_{t+1})) - q(a_t|s_t))$$

$$q(a_t|s_t) = (1 - \alpha)q(a_t|s_t) + \alpha(R_t + \gamma q(a_{t+1}|s_{t+1}))$$

POLICY UPDATE

ϵ – greedy approach

SARSA

$$q(a_t|s_t) = (1 - \alpha)q(a_t|s_t) + \alpha(R_t + \gamma q(a_{t+1}|s_{t+1}))$$

Q LEARNING

REQUIRES ONLY SARS PAIRS

$$q(a_t|s_t) = (1 - \alpha)q(a_t|s_t) + \alpha * \max(R_t + \gamma q(a_*|s_{t+1}))$$

DEMO TIME

(GRID-WORLD)

GOAL OF Q LEARNING

- ESTIMATING ACTION VALUE FUNCTION (Q)
- UPDATE Q FUNCTION VIA MOVING AVG

$$q(a_t|s_t) = (1 - \alpha)q(a_t|s_t) + \alpha * \max(R_t + \gamma q(a_*|s_{t+1}))$$


PROBLEM WITH Q LEARNING

- CAN'T HANDLE NEAR INFINITE STATES (IMAGES)
- CAN'T HANDLE **S** STATE AS FEATURE VECTOR

Q LEARNING

$$q(a_t|s_t) = (1 - \alpha)q(a_t|s_t) + \alpha * \max(R_t + \gamma q(a_*|s_{t+1}))$$


DEEP Q LEARNING




Update neural network by setting label for $q(a_t|s_t)$ as $\max(R_t + \gamma q(a_*|s_{t+1}))$

$$q(a_t|s_t) = \text{neural} - \text{network}(s_t)$$

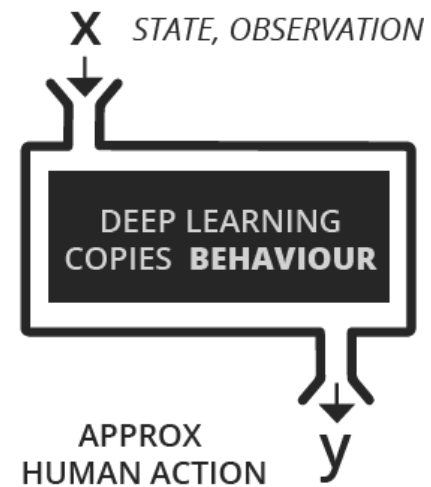
FUNCTION APPROXIMATOR



DEEP Q LEARNING



BEHAVIORAL CLONING

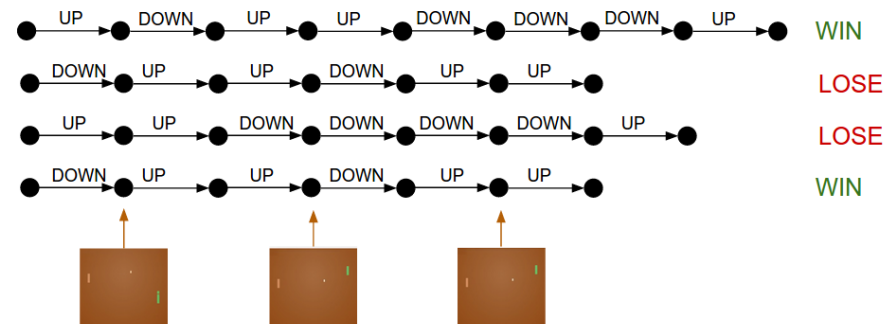



DEMO TIME

(CARTPOLE IN OPENAI-GYM)

POLICY GRADIENT APPROACHES

CHANGE POLICY RATHER THAN Q VALUES





POLICY ENTROPY = 0.9



POLICY ENTROPY = 0.1

ACTOR CRITIC MODEL




ACTOR

*Learns to act based on
advantage factor
given by critic.*

CRITIC

*Understands each shot(State),
based on rewards*

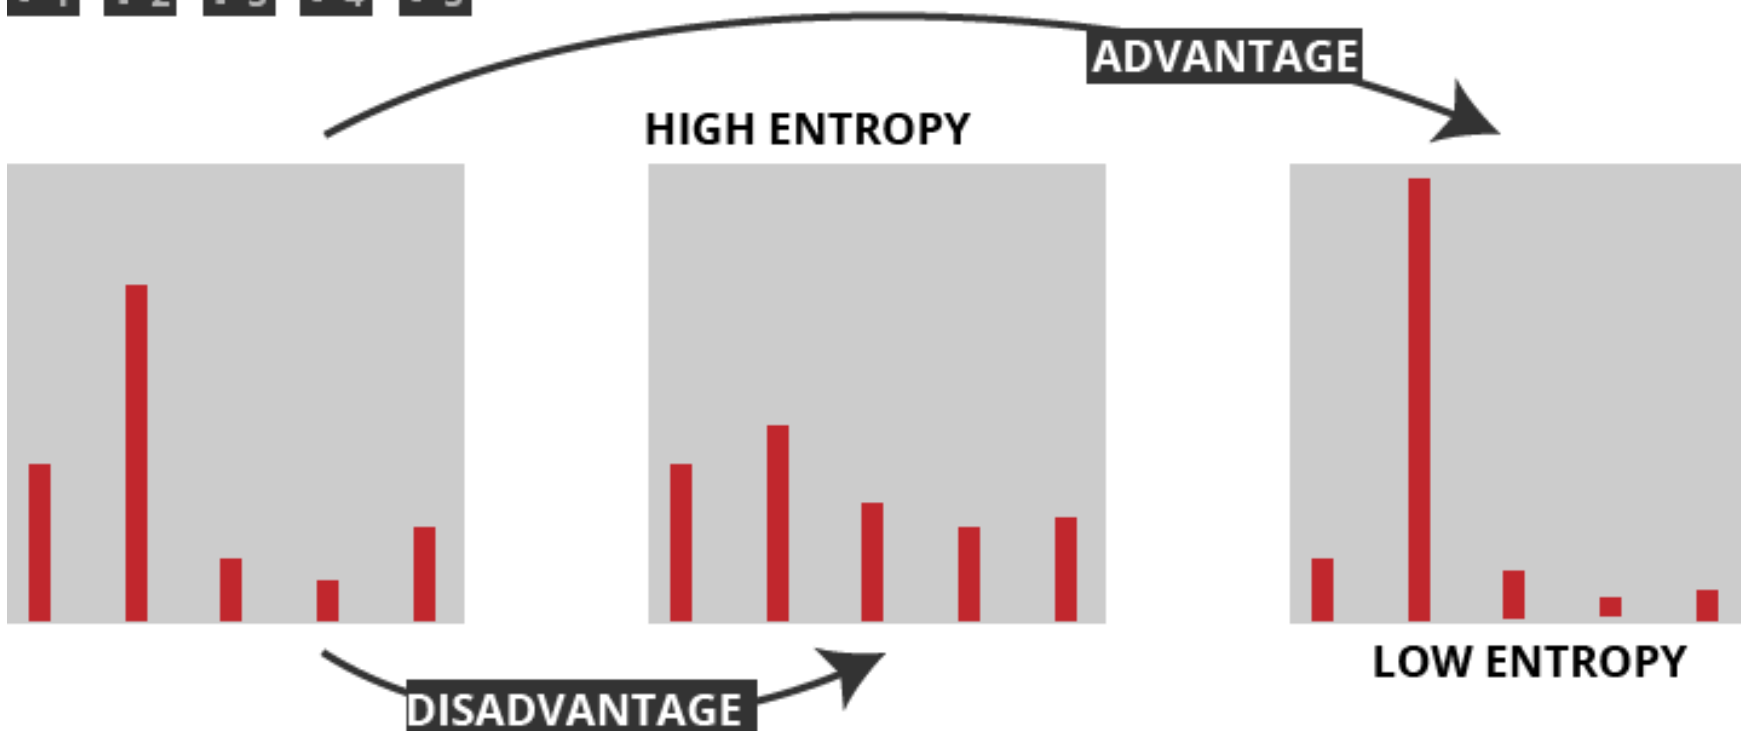




$$\begin{aligned}
 & \text{— EXPECTED REWARD} \\
 & \text{— PREDICTED REWARD} \\
 & \text{= ADVANTAGE}
 \end{aligned}$$

$$R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-t-2} R_T$$

p_1 p_2 p_3 p_4 p_5




OTHER COOL STUFF

PROXIMAL POLICY OPTIMIZATION
MONTE CARLO TREE SEARCH

Subscribe

www.youtube.com/c/crazymuse



Contact Details

jaley.dholakiya@crazymuse.in
8105207901

Github link : www.github.com/crazymuse

Medium Link : www.medium.com/crazymuse

Ping me for any doubt